




UNIZG-FER 86518
Service-Oriented Computing



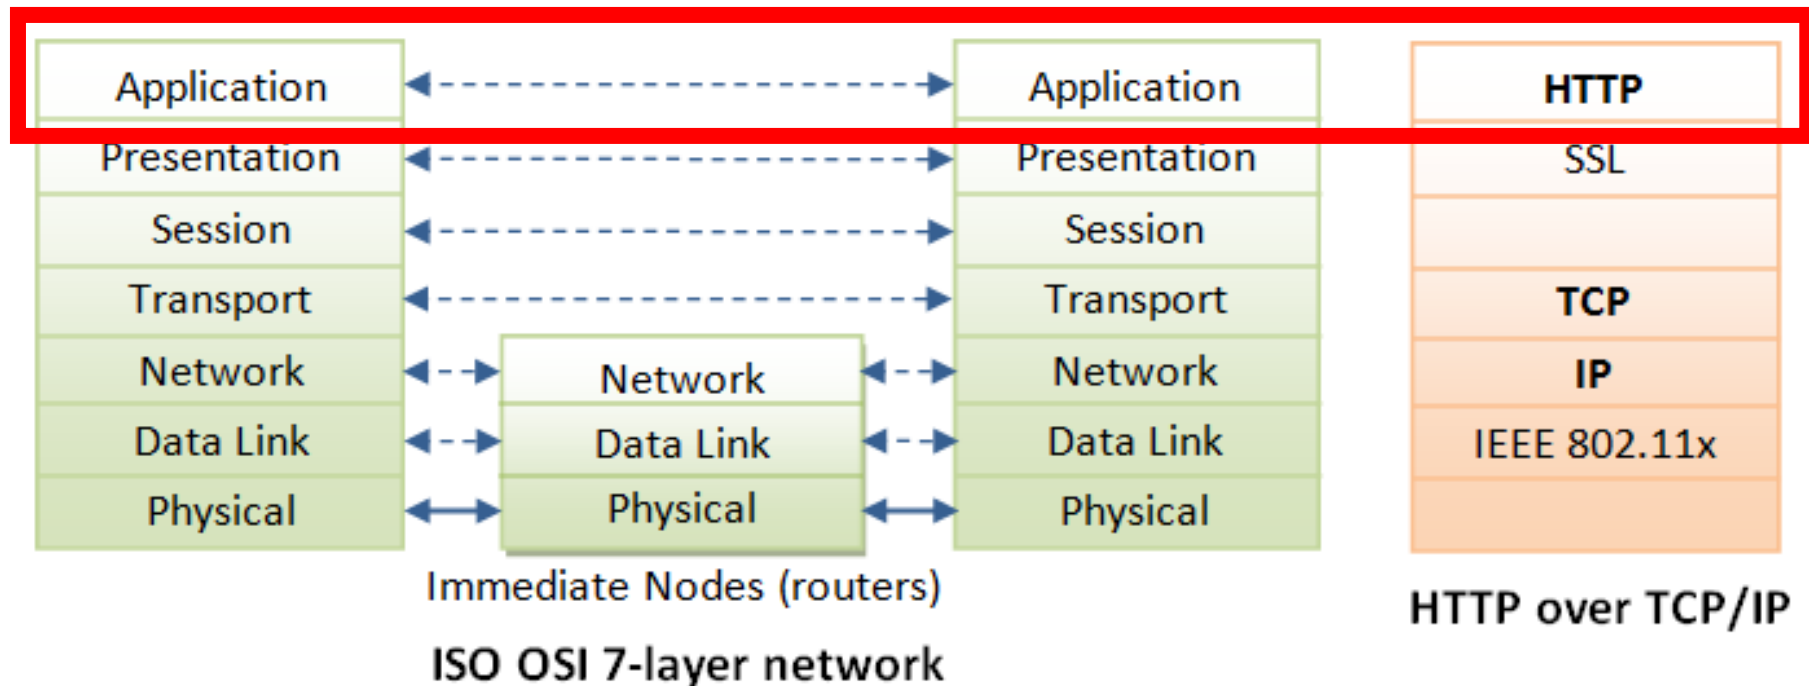
HTTP

The Driver of the World Wide Web

HTTP Basics





- HTTP (*HyperText Transfer Protocol*) protocol
 - Application-level protocol for distributed hypermedia information systems

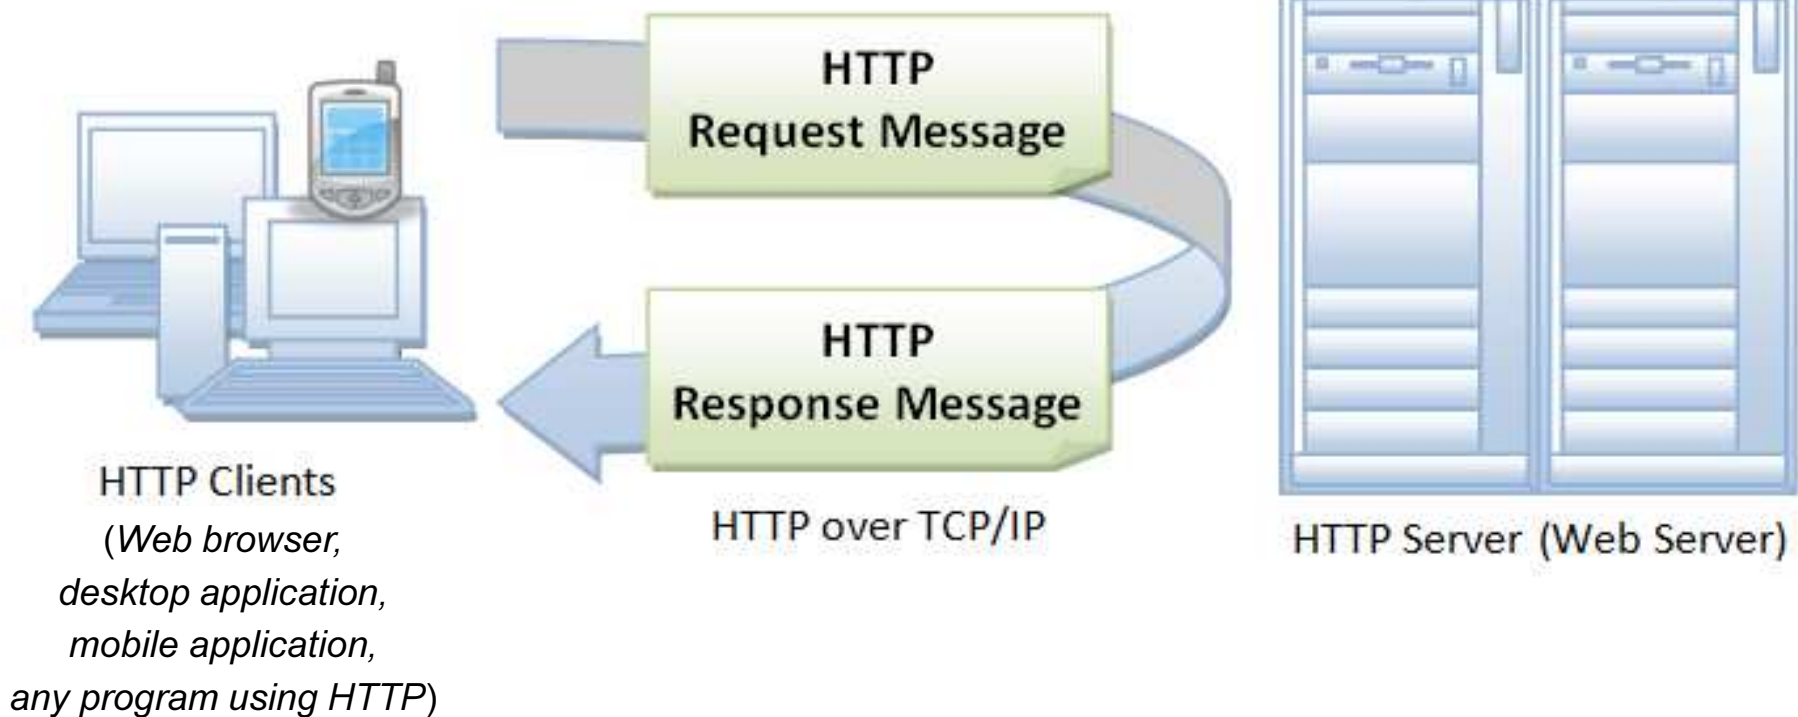


- **World Wide Web**
 - Applications and services based on **HTTP protocol**

HTTP Basics


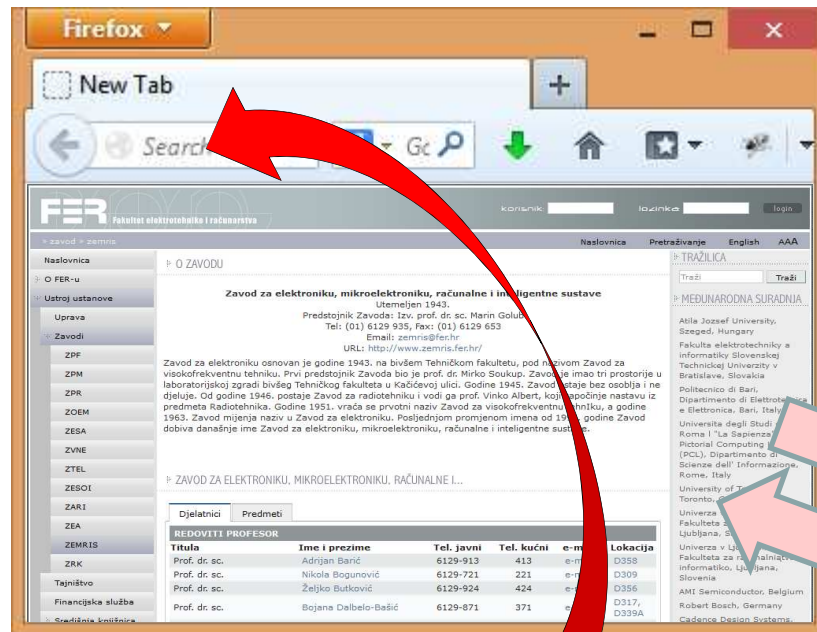


- HTTP (*HyperText Transfer Protocol*) protocol
 - Client-server architecture
 - Asymmetric request-response protocol
 - Client *pulls* information from the server
(instead of server *pushes* information down to the client)



HTTP Basics

- Web browser
 - Most common HTTP client

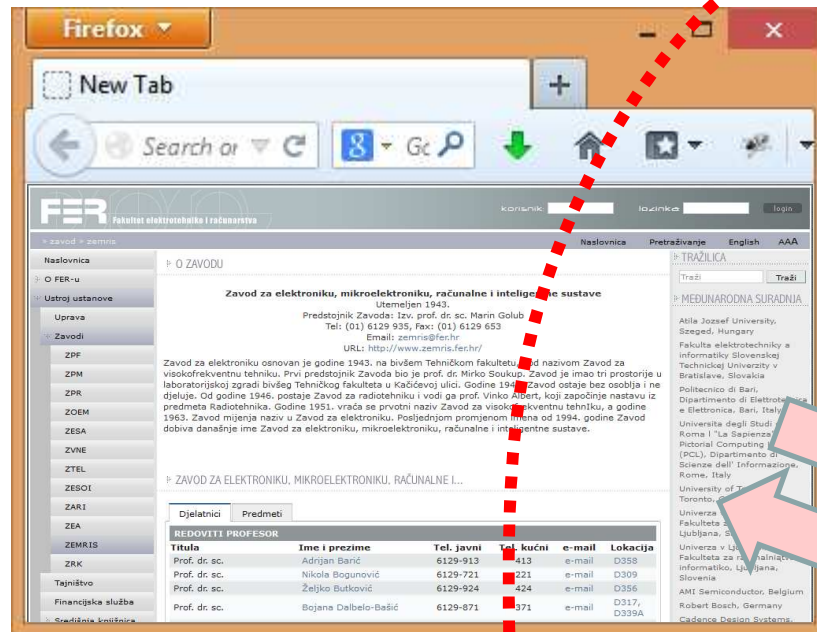


`http://www.fer.unizg.hr/zavod/zemris`

protocol *host* *resource path*

HTTP Basics

- Web browser
 - Most common HTTP client



`http://www.fer.unizg.hr/zavod/zemris`

`http`
protocol

`www.fer.unizg.hr`
host

`/zavod/zemris`
resource path

DNS server
208.43.0.4

IP of `www.fer.unizg.hr`?

161.53.72.120

Network connection (socket)



HTTP request: `GET /zavod/zemris`

HTTP response



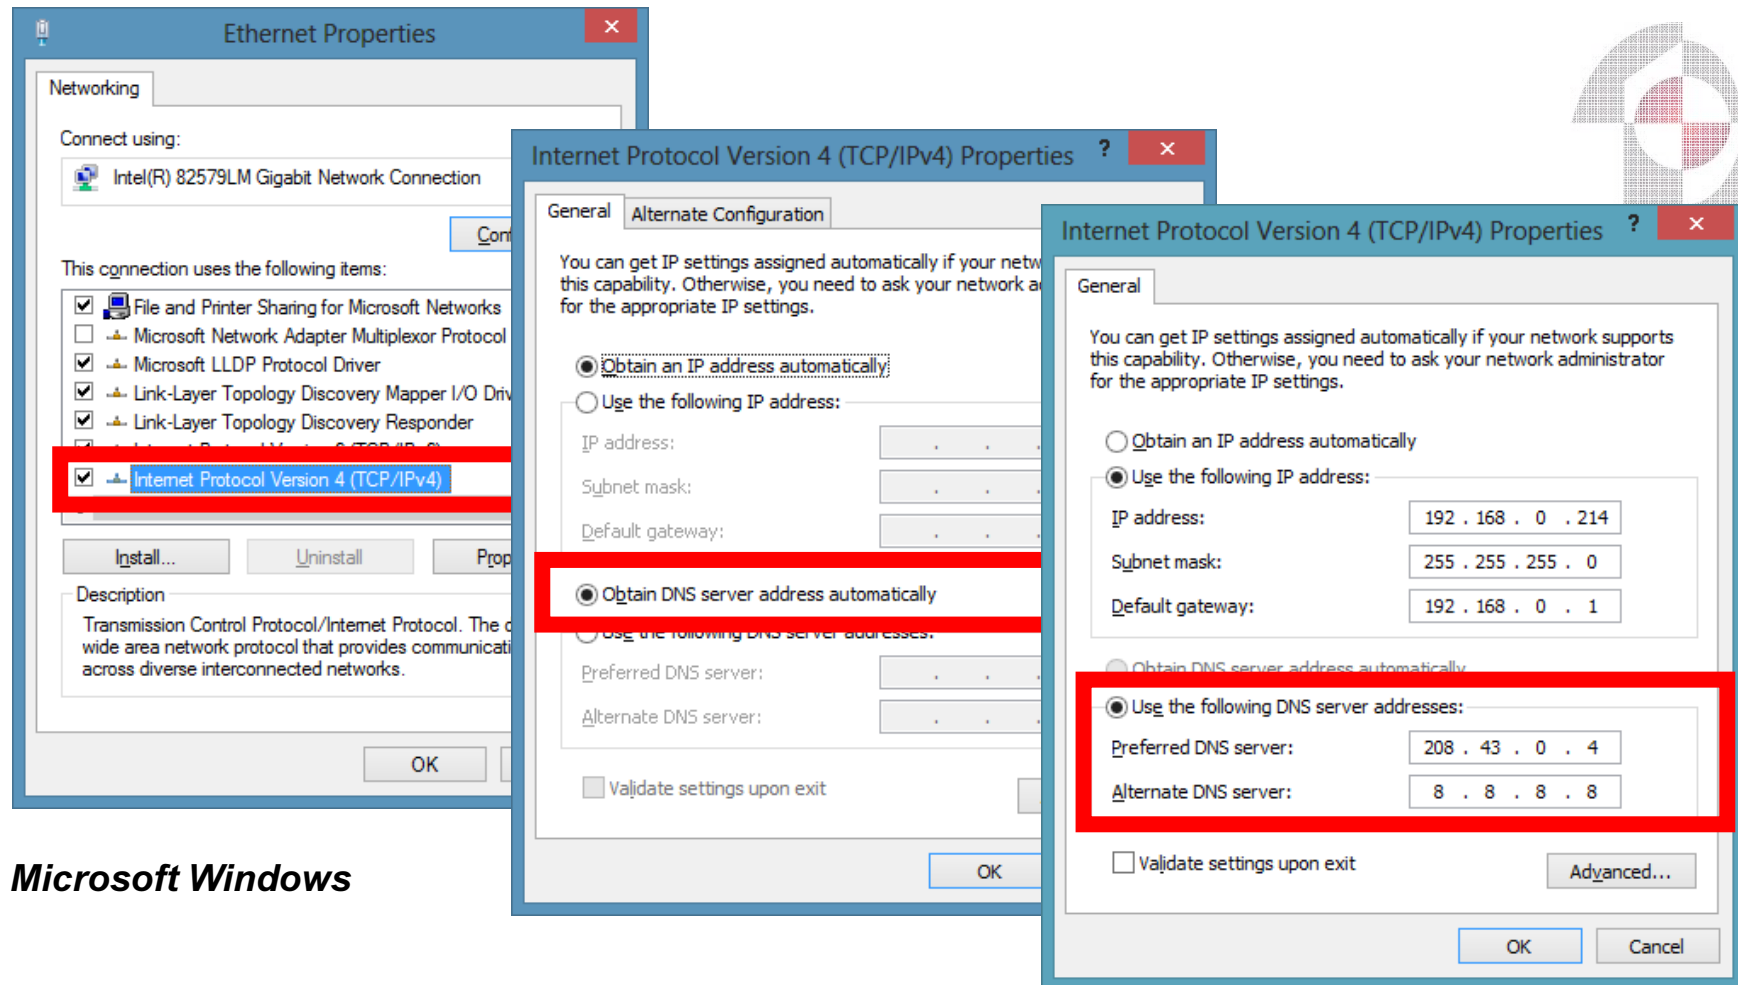
Web server

`www.fer.unizg.hr`
161.53.72.120



HTTP Basics

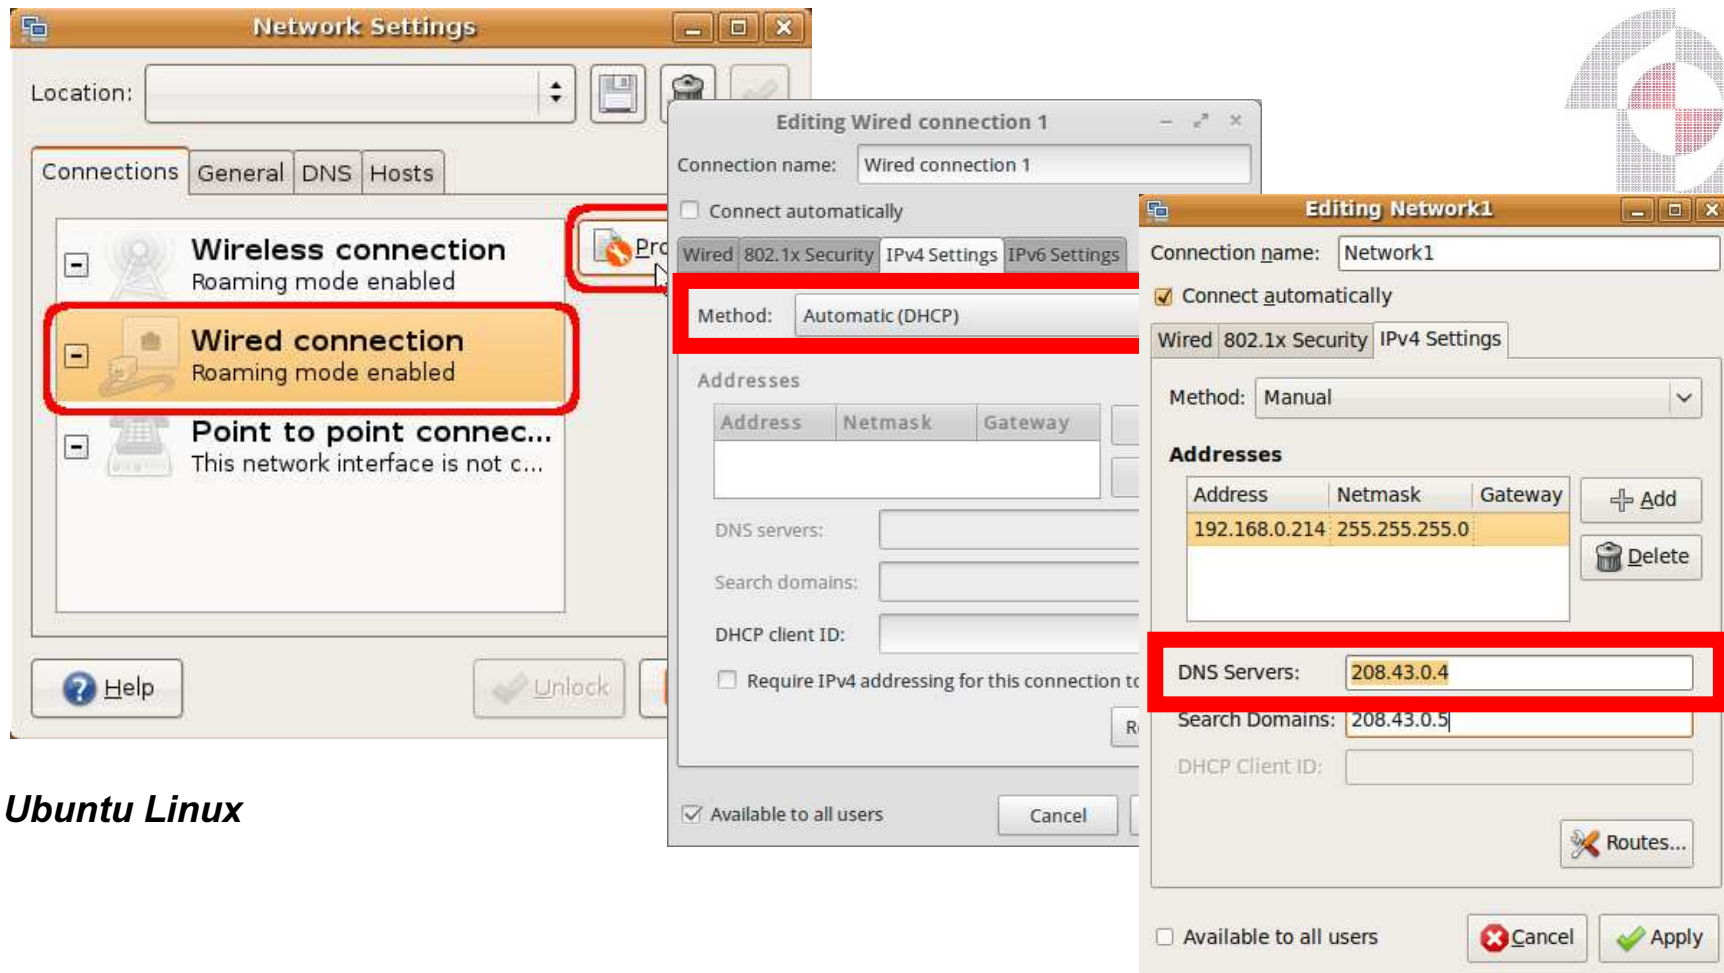
- How web browser finds the DNS server?
 - Operating system network configuration



Microsoft Windows

HTTP Basics

- How web browser finds the DNS server?
 - Operating system network configuration



Ubuntu Linux

HTTP Basics

- Uniform Resource Identifier (URI)
 - String used to uniquely identify a resource over the web
- URI syntax

```
protocol://hostname:port/path-and-file-name?parameters
```

protocol

- The application-level protocol used by the client and server
e.g. HTTP, HTTPS, FTP, telnet

hostname

- The DNS domain name or IP address of the server
e.g. www.fer.unizg.hr, 161.53.72.120

port



- The TCP port number the server is listening for incoming requests from the clients

path-and-file-name

- The name and location of the requested resource under the server document base directory
e.g. static file on disk or program that dynamically renders the response

parameters

- Optional, used to additionally describe the resource (*we'll come back to this later*)



HTTP Basics

- URI examples

1) `http://www.fer.unizg.hr/zavod/zemris` (default HTTP port is 80)

`http://www.fer.unizg.hr:80/zavod/zemris`

`http://161.53.72.120/zavod/zemris`

`http://161.53.72.120:80/zavod/zemris`

2) `http://www.example.com:1234/europe/croatia/home.html`



3) `https://www.fer.unizg.hr/predmet/rznu` (default HTTPS port is 443)

`https://www.fer.unizg.hr:443/predmet/rznu`



4) `https://www.fer.unizg.hr:987/predmet/rznu`

5) `ftp://www.ftp.org/docs/test.txt` (default FTP port is 21)

6) `telnet://www.test101.com/` (default TELNET port is 23)



HTTP Basics




General HTTP client algorithm

1. The user enters URI of a desired web page
 2. Browser parses the URI
 3. Browser asks the DNS server for web server's IP address
 4. DNS server responds with IP address
- (steps 3 and 4 are not necessary if user enters web server's IP address instead of DNS name)*
5. Browser opens a network connection to given IP address and TCP port
 6. Browser sends a HTTP request message to the web server
 7. Server maps the URI to a local file or program
 8. Server returns a HTTP response message
 9. Browser formats the response, renders GUI, and displays a web page

HTTP Basics

- Other HTTP clients

- curl



- Command line syntax



`curl [options] <url>`

- Usage instructions

`curl -help`



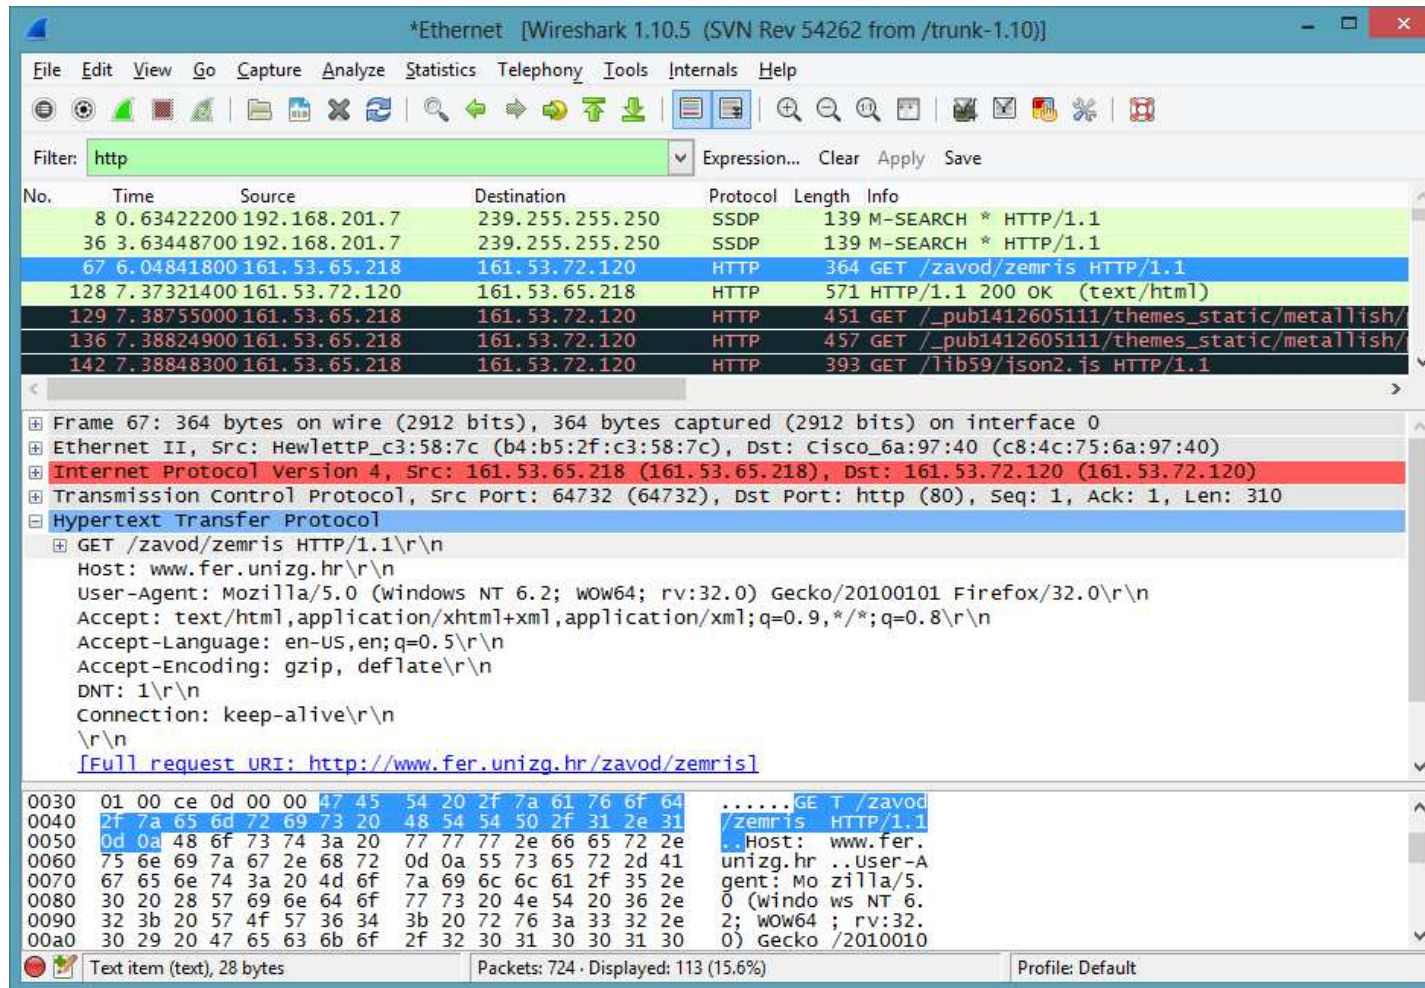
- Simple HTTP request

`curl http://www.fer.unizg.hr/zavod/zemris`




HTTP Basics

- What happens on the network level?
 - Network monitoring & capture tool



HTTP Basics



- What happens on the network level?

HTTP request message


Web browser (Firefox)

```
GET /zavod/zemris HTTP/1.1
Host: www.fer.unizg.hr
User-Agent: Mozilla/5.0 (Windows NT 6.2; WOW64; rv:23.0) Gecko/20100101 Firefox/23.0
Accept: text/html,application/xhtml+xml,application/xml;q=0.9,*/*;q=0.8
Accept-Language: en-US,en;q=0.5
Accept-Encoding: gzip, deflate
DNT: 1
Cookie: __utma=161902635.17065694.1374136092.1377600687.1377615217.13; __utmz=...
Connection: keep-alive
```

curl

```
GET /zavod/zemris HTTP/1.1
User-Agent: curl/7.32.0
Host: www.fer.unizg.hr
Accept: */*
```

HTTP Basics



- What happens on the network level?



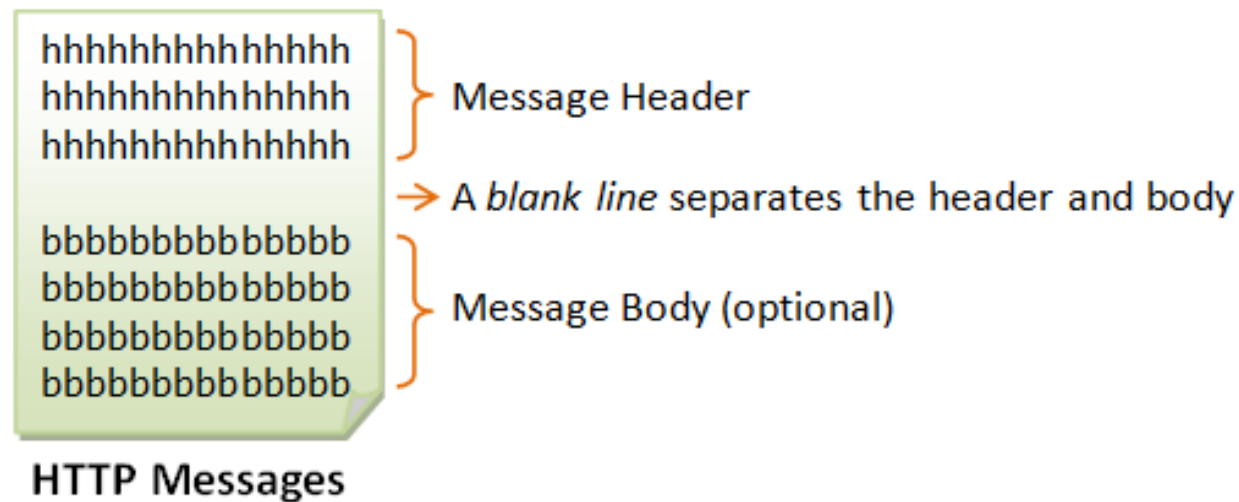
HTTP response message

```
HTTP/1.1 200 OK
Date: Wed, 28 Aug 2013 10:49:28 GMT
Server: Apache/2.2.23 (FreeBSD) mod_fcgid/2.3.6 mod_ssl/2.2.23 OpenSSL/0.9.8x
X-Powered-By: PHP/5.3.19
Expires: Thu, 19 Nov 1981 08:52:00 GMT
Cache-Control: no-store, no-cache, must-revalidate, post-check=0, pre-check=0
Pragma: no-cache
P3P: CP="NOI CURa ADMa DEVa TAIa PSAa PSDa IVAa IVDa HISa OTPa ..."
Set-Cookie: CMS=vhtenqteedgso8d9u6l8h8vgq6; expires=Wed, 04-Sep-2013...
Content-Length: 73120
Content-Type: text/html; charset=utf-8

<!DOCTYPE html PUBLIC "-//W3C//DTD XHTML 1.0 Transitional//EN"
"http://www.w3.org/TR/xhtml1/DTD/xhtml1-transitional.dtd">
<html xmlns="http://www.w3.org/1999/xhtml" lang="hr"><head><meta http-equiv="Content-
Type" content="text/html; charset=utf-8" /><meta http-equiv="Content-Language"
content="hr" /><meta name="generator" content="QuiltCMS 2.0, http://www.fer.hr/" />
<!--meta name="robots" content="noindex" /--><meta name="keywords" content=",, />
<title>Zavod za elektroniku, mikroelektroniku, računalne i inteligentne sustave - FER
e-Campus v1</title><base href="http://www.fer.unizg.hr/zavod/zemris"><link
rel="alternate" type="application/rss+xml" title="FER: Zavod za elektroniku,
mikroelektroniku...
```

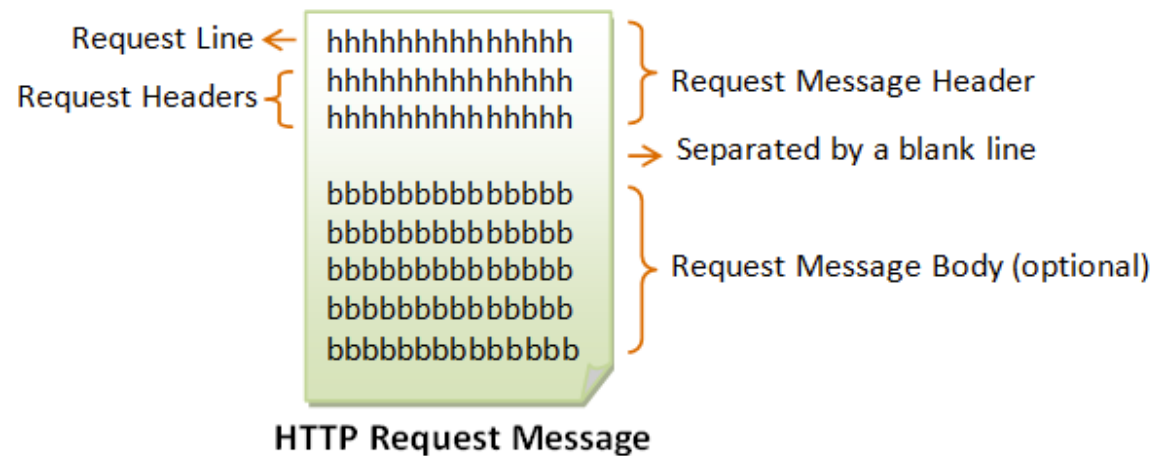

HTTP Basics

- HTTP messages
 - General form
 - Each HTTP message (either request or response) follows this general form



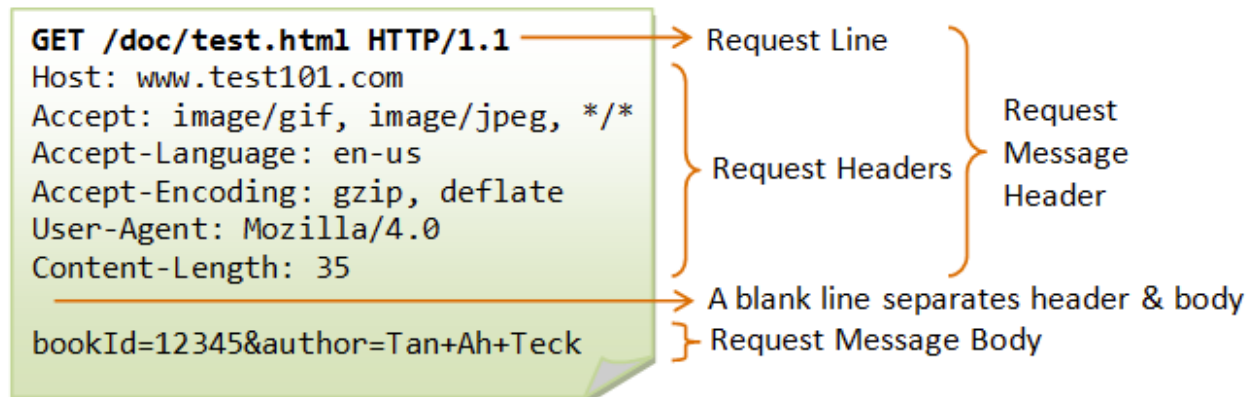


HTTP Basics

- HTTP request message




Example:



HTTP Basics

- HTTP request message



Request line:

`request-method-name request-URI HTTP-version CRLF`

request-method-name

Informs the server which operation to perform over the resource

HTTP protocol defines a set of request methods: GET, PUT, POST, DELETE, HEAD, and OPTIONS

The client uses one of these methods to send a request to the server



request-URI

Specifies the resource on a web server over which the server should perform the requested operation

HTTP-version


Client specifies the version of HTTP protocol it understands

Two versions are currently in use: HTTP/1.0 and HTTP/1.1



HTTP Basics

- HTTP request message





Request line:

request-method-name request-URI HTTP-version CRLF


Examples:

```
GET /zavod/zemris HTTP/1.1
GET /zavod/zemris HTTP/1.0
HEAD /zavod/zemris HTTP/1.1
PUT /photoalbum/image03.jpg HTTP/1.0
DELETE /photoalbum/image03.jpg HTTP/1.0
POST /news/article/comments HTTP/1.1
```



HTTP Basics



- HTTP request message



Request header:

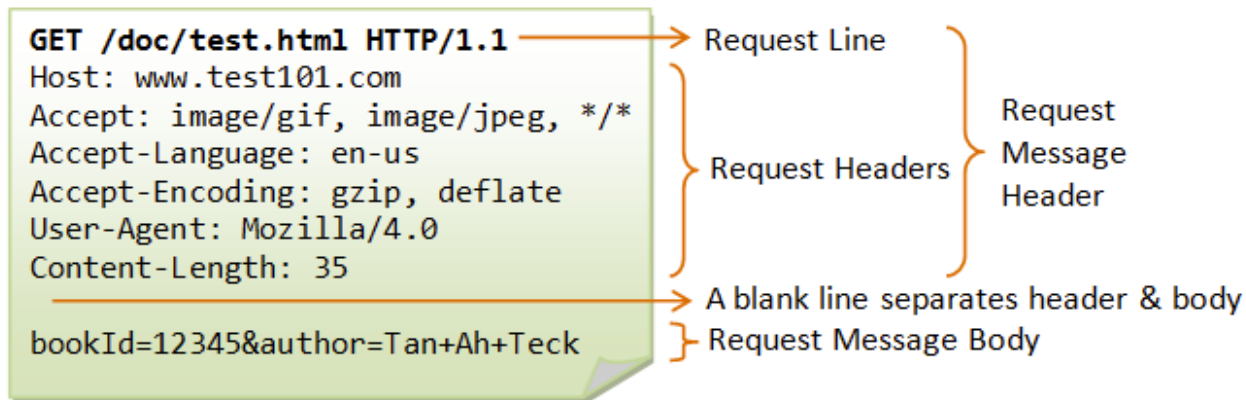
`request-header-name: request-header-value1, request-header-value2, ... CRLF`

- The request headers are in the form of `name: value` pairs
- Multiple header values, separated by commas, can be specified
- Each request header ends with a new line (CRLF)
- HTTP allows arbitrary number of request headers in single request
- HTTP also allows custom non-standard header names
(*custom web servers might process custom headers, standard web servers ignore them*)



HTTP Basics



- HTTP request message



Request message body:

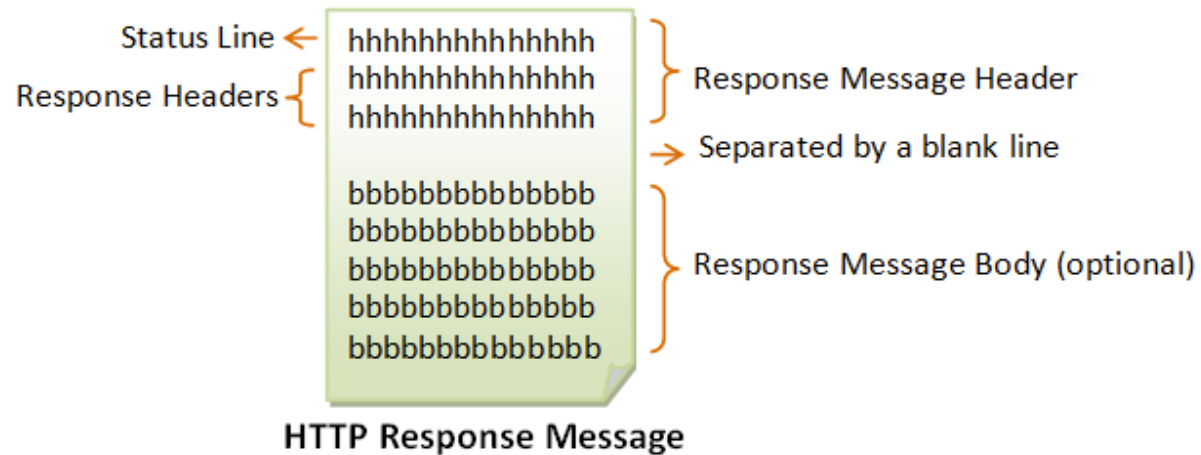
no defined structure, free format, arbitrary length

- Optional part of HTTP request message
- Used to send extra data with the request that cannot be specified in request headers
(for example, user-defined parameters)
- HTTP protocol does not define the structure of request message body
- HTTP headers specify how to interpret the body
(for example, Content-Type header)



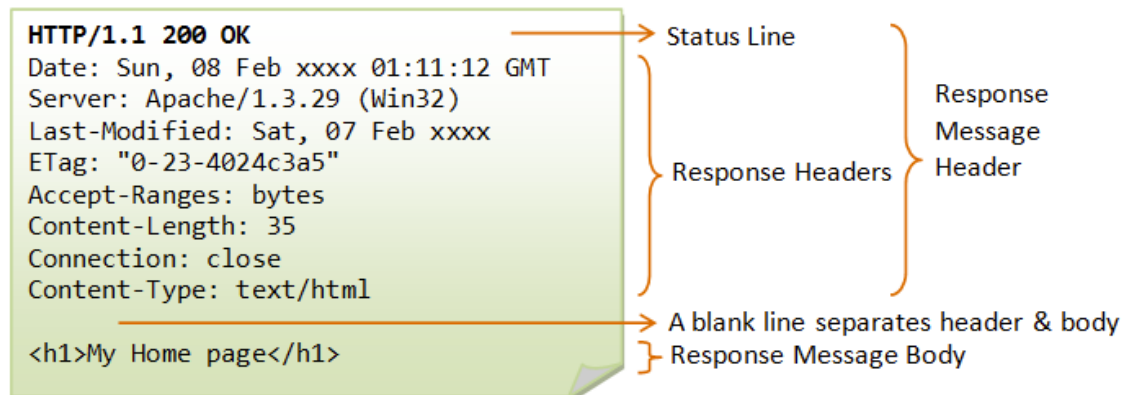


HTTP Basics

- HTTP response message




Example:



HTTP Basics

- HTTP response message



Status line:

HTTP-version status-code reason-phrase CRLF

HTTP-version

Server specifies the version of the HTTP protocol used in response

Version chosen by server should be equal or lower than the version specified in client's request

Two versions are currently in use: HTTP/1.0 and HTTP/1.1



status-code

A 3-digit number generated by the server to reflect the outcome of the request

Inform the client whether request is served successfully, some error occurred, etc.

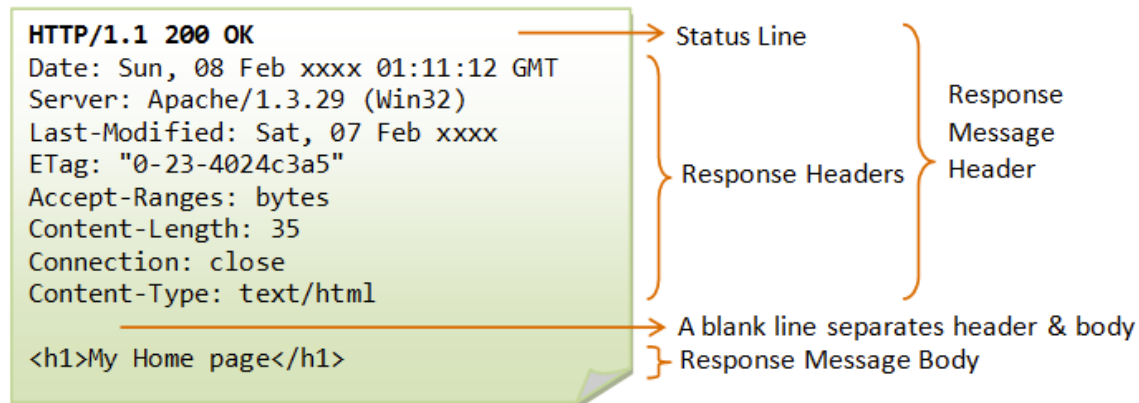
reason-phrase

Gives a short explanation to the status code



HTTP Basics

- HTTP response message



Status line:

HTTP-version status-code reason-phrase CRLF



Examples:

HTTP/1.1 200 OK

HTTP/1.0 404 Not Found

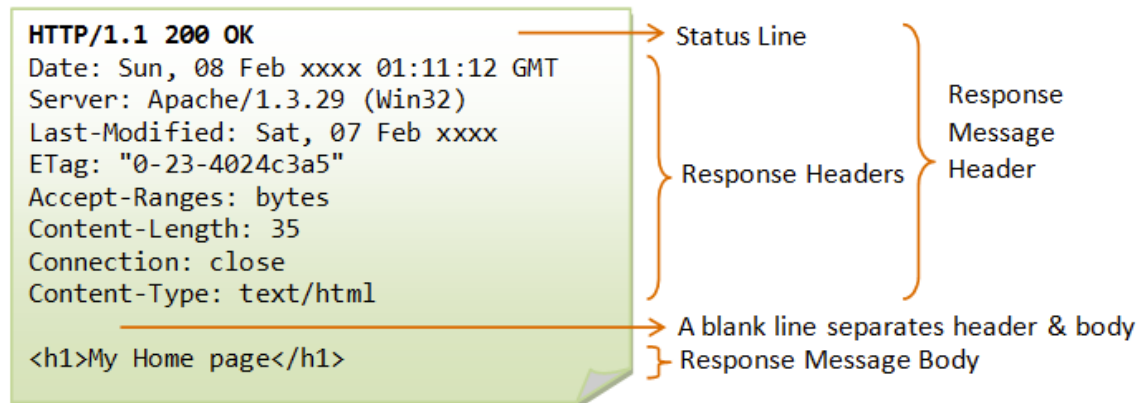
HTTP/1.1 403 Forbidden

HTTP/1.1 500 Internal Server Error



HTTP Basics



- HTTP response message



Response header:


`response-header-name: resp-header-value1, resp-header-value2, ... CRLF`

- Response headers follow the same form as request headers
 - The response headers are in the form of `name: value` pairs
 - Multiple header values, separated by commas, can be specified
 - Each response header ends with a new line (CRLF)
 - HTTP allows arbitrary number of response headers in single request
 - HTTP also allows custom non-standard header names
(*custom clients might process custom headers, standard HTTP clients ignore them*)



HTTP Basics



- HTTP response message




Response message body:

no defined structure, free format, arbitrary length

- Optional part of HTTP response message
- Used to send data from web server back to the client
(for example, web page's HTML, client-side script, image)
- HTTP protocol does not define the structure of response message body
- HTTP response headers specify how to interpret the body
(for example, Content-Type header)



HTTP Basics



- HTTP client socket-level programming

```
import java.net.*;
import java.io.*;

public class HttpClientSocket {
    public static void main(String[] args) throws IOException {
        // The host and port to be connected
        String host = "www.fer.unizg.hr";
        int port = 80;
        // Create a TCP socket and connect to the host:port
        Socket socket = new Socket(host, port);

        // Create the input and output streams for the network socket
        BufferedReader in = new BufferedReader(new InputStreamReader(socket.getInputStream()));
        PrintWriter out = new PrintWriter(socket.getOutputStream(), true);

        // Create request line
        out.println("GET /zavod/zemris HTTP/1.1");
        // Add some request headers
        out.println("Host: www.unizg.fer.hr");
        out.println("User-Agent: My custom HTTP client");
        // Add blank line separating header & body
        out.println();
        // Send request to the HTTP server
        out.flush();

        // Read the response and display on console
        String line;
        // readLine() returns null if server close the network socket
        while((line = in.readLine()) != null) {
            System.out.println(line);
        }
        // Close the I/O streams
        in.close();
        out.close();
    }
}
```


HTTP Basics



- HTTP client socket-level programming

- Compile program


```
javac HttpClientSocket.java
```

- Start program

```
java HttpClientSocket
```



HTTP Basics



- HTTP client programming using HTTP library

```
import java.net.*;
import java.io.*;

public class HttpClientHttpLib {
    public static void main(String[] args) throws IOException {
        // The URI of the remote resource
        String uri = "http://www.fer.unizg.hr/zavod/zemris";
        // Open a TCP connection for HTTP communication with a resource with given URI
        HttpURLConnection http = (HttpURLConnection) new URL(uri).openConnection();

        // Read response status
        System.out.println("Response status code: " + http.getResponseCode());
        System.out.println("Response reason phrase: " + http.getResponseMessage());

        // Read response data if any
        String line;
        BufferedReader in = new BufferedReader(new InputStreamReader(http.getInputStream()));
        while((line = in.readLine()) != null) {
            System.out.println(line);
        }

        // Close the connection
        http.disconnect();
    }
}
```

HTTP Basics

- HTTP client programming using HTTP library

- Compile program

```
javac HttpClientHttpLib.java
```

- Start program

```
java HttpClientHttpLib
```

